

Network Working Group
Internet-Draft
Intended status: Informational
Expires: March 31, 2019

L. Qiang, Ed.
B. Liu
T. Eckert, Ed.
Huawei
L. Geng
L. Wang
China Mobile
September 27, 2018

Large-Scale Deterministic Network
draft-qi-ang-detnet-large-scale-detnet-02

Abstract

This document presents the framework and key methods for Large-scale Deterministic Networks (LDN). It achieves scalability for the number of supportable deterministic traffic flows via Scalable Deterministic Forwarding (SDF) that does not require per-flow state in transit nodes and precise time synchronization among nodes. It achieves Scalable Resource Reservation (SRR) by allowing for it to be decoupled from the forwarding plane nodes, and aggregating resource reservation status in time slots.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of [BCP 78](#) and [BCP 79](#).

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on March 31, 2019.

Copyright Notice

Copyright (c) 2018 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to [BCP 78](#) and the IETF Trust's Legal Provisions Relating to IETF Documents

(<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	2
1.1. Requirements Language	3
1.2. Terminology & Abbreviations	3
2. Overview	4
2.1. Summary	4
2.2. Background	4
2.2.1. Deterministic End-to-End Latency	4
2.2.2. Hop-by-Hop Delay	4
2.2.3. Cyclic Forwarding	5
2.2.4. Co-Existence with Non-Deterministic Traffic	5
2.3. System Components	6
3. Scalable Deterministic Forwarding	7
3.1. Three Queues	8
3.2. Cycle Mapping	9
3.2.1. Cycle Identifier Carrying	9
4. Scalable Resource Reservation	10
5. Performance Analysis	11
5.1. Queueing Delay	11
5.2. Jitter	12
6. IANA Considerations	14
7. Security Considerations	14
8. Acknowledgements	14
9. Normative References	14
Authors' Addresses	15

1. Introduction

Deploying deterministic service over large-scale network will face some technical challenges, such as

- o massive number of deterministic flows vs. per-flow operation and management;
- o long link propagation may bring in significant jitter;
- o time synchronization is hard to be achieved among numerous devices, etc.

Motivated by these challenges, this document presents a Large-scale Deterministic Network (LDN) system, which consists of Scalable Deterministic Forwarding (SDF) at forwarding plane and Scalable Resource Reservation (SRR) at control plane. The technologies of SDF and SRR can be used independently.

As [[draft-ietf-detnet-problem-statement](#)] indicates, deterministic forwarding can only apply on flows with well-defined traffic characteristics. The traffic characteristics of DetNet flow has been discussed in [[draft-ietf-detnet-architecture](#)], that could be achieved through shaping at Ingress node or up-front commitment by application. This document assumes that DetNet flows follow some specific traffic patterns accordingly.

1.1. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC 2119](#).

1.2. Terminology & Abbreviations

This document uses the terminology defined in [[draft-ietf-detnet-architecture](#)].

TSN: Time Sensitive Network

CQF: Cyclic Queuing and Forwarding

LDN: Large-scale Deterministic Network

SDF: Scalable Deterministic Forwarding

SRR: Scalable Resource Reservation

DSCP: Differentiated Services Code Point

EXP: Experimental

TC: Traffic Class

T: the length of a cycle

H: the number of hops

K: the size of aggregated resource reservation window

2. Overview

2.1. Summary

The Large-Scale Deterministic Network solution (LDN) consists of two parts: The Scalable Deterministic Forwarding Plane (SDF) as its forwarding plane and the Scalable Resource Reservation (SRR) as its control plane. In the SDF, nodes in the network have synchronized frequency, and each node forwards packets in a slotted fashion based on a cycle identifiers carried in packets. Ingress nodes or senders have a function called gate to shape/condition traffic flows. Except for this gate function, the SDF has no awareness of individual flows. The SRR maintains resource reservation states for deterministic flows, Ingress nodes maintain per-flow states and core nodes aggregate per-flow states in time slots.

2.2. Background

This section motivates the design choices taken by the proposed solution and gives the necessary background for deterministic delay based forwarding plane designs.

2.2.1. Deterministic End-to-End Latency

Bounded delay is delay that has a deterministic upper and lower bound.

The delay for packets that need to be forwarded with deterministic delay needs to be deterministic on every hop. If any hop in the network introduces non-deterministic delay, then the network itself can not deliver a deterministic delay service anymore.

2.2.2. Hop-by-Hop Delay

Consider a simple example (without picture), where N has 10 receiving interfaces and one outgoing interface I all of the same speed. There are 10 deterministic traffic flows, each consuming 5% of a links bandwidth, one from each receiving interface to the outgoing interface.

Node N sends 'only' 50% deterministic traffic to interface I, so there is no ongoing congestion, but there is added delay. If the arrival time of packets for these 10 flows into N is uncontrolled, then the worst case is for them to all arrive at the same time. One packet has to wait in N until the other 9 packets are sent out on I, resulting in a worst case deterministic delay of 9 packets serialization time. On the next hop node N2 downstream from N, this problem can become worse. Assume N2 has 10 upstream nodes like N,

the worst case simultaneous burst of packets is now 100 packets, or a 99 packet serialization delay as the worst case upper bounded delay incurred on this hop.

To avoid the problem of high upper bound end-to-end delay, traffic needs to be conditioned/interleaved on every hop. This allows to create solutions where the per-hop-delay is bounded purely by the physics of the forwarding plane across the node, but not the accumulated characteristics of prior hop traffic profiles.

2.2.3. Cyclic Forwarding

The common approach to solve that problem is that of a cyclic hop-by-hop forwarding mechanism. Assume packets forwarded from N1 via N2 to N3 as shown in Figure 1. When N1 sends a packet P to interface I1 with a Cycle X, it must be guaranteed by the forwarding mechanism that N2 will forward P via I2 to N3 in a cycle Y.

The cycle of a packet can either be deduced by a receiving node from the exact time it was received as is done in SDN/TDMA systems, and/or it can be indicated in the packet. This document solution relies on such markings because they allow to reduce the need for synchronous hop-by-hop transmission timings of packets.

In a packet marking based slotted forwarding model, node N1 needs to send packets for cycle X before the latest possible time that will allow for N2 to further forward it in cycle Y to N3. Because of the marking, N1 could even transmit packets for cycle X before all packets for the previous cycle (X-1) have been sent, reducing the synchronization requirements between across nodes.




Figure 1: Cyclic Forwarding

2.2.4. Co-Existence with Non-Deterministic Traffic

Traffic with deterministic delay requirements can co-exist with traffic only requiring non-deterministic delay by using packet scheduling where the delay incurred by non-deterministic packets is deterministic for the deterministic traffic (and low). If LDN SDF is deployed together with such non-deterministic delay traffic than such a scheme must be supported by the forwarding plane. A simple approach for the delay incurred on the sending interface of a

deterministic node due to non-deterministic traffic is to serve deterministic traffic via a strict, highest-priority queue and include the worst case delay of a currently serialized non-deterministic packet into the deterministic delay budget of the node. Similar considerations apply to the internal processing delays in a node.

2.3. System Components

The Figure 2 shows an overview of the components considered in this document system and how they interact.

A network topology of nodes, Ingress, Core and Egress support a method for cyclic forwarding to enable Scalable Deterministic Forwarding (SDF). This forwarding requires no per-flow state on the nodes.

Ingress edge nodes may support the (G)ate function to shape traffic from sources into the desired traffic characteristics, unless the source itself has such function. Per-flow state is required on the ingress edge node.

A Scalable Resource Reservation (SRR) works as control plane. It records reserved resources for deterministic flows. Per-flow state is maintained on the ingress edge node, and aggregated state is maintained on core node.

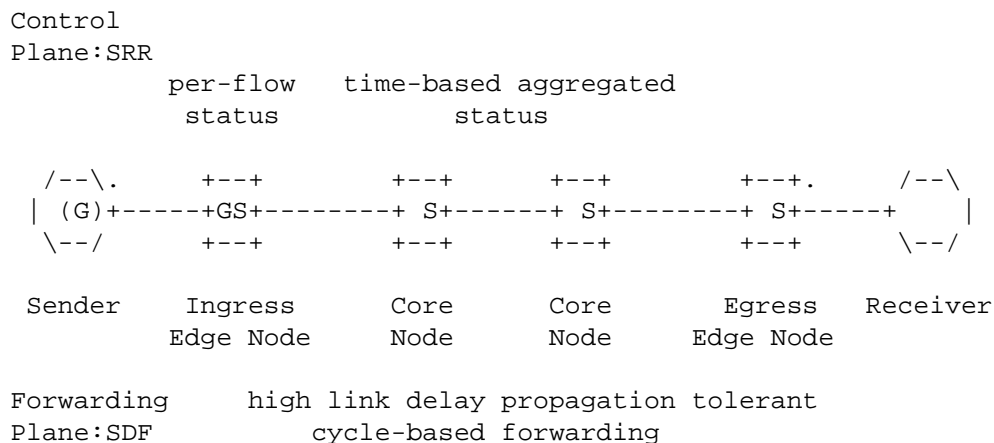


Figure 2: System Overview

3. Scalable Deterministic Forwarding

DetNet aims at providing deterministic service over large scale network. In such large scale network, it is difficulty to get precise time synchronization among numerous devices. To reduce requirements, the forwarding mechanism described in this document assumes only frequency synchronization but not time synchronization across nodes: nodes maintain the same clock frequency $1/T$, but do not require the same time as shown in Figure 3.




Figure 3: Time Synchronization & Clock Synchronization

IEEE 802.1 CQF is an efficient forwarding mechanism in TSN that guarantees bounded end-to-end latency. CQF is designed for limited scale networks. Time synchronization is required, and the link propagation delay is required to be smaller than a cycle length T . Considering the large scale network deployment, the proposed Scalable Deterministic Forwarding (SDF) permits frequency synchronization and link propagation delay may exceed T . Besides these two points, CQF and the asynchronous forwarding of SDF are very similar.

Figure 4 compares CQF and SDF through an example. Suppose Node A is the upstream node of Node B. In CQF, packets sent from Node A at cycle x , will be received by Node B at the same cycle, then further be sent to downstream node by Node B at cycle $x+1$. Due to long link propagation delay and frequency synchronization, Node B will receive packets from Node A at different cycle denoted by y in the SDF, and Node B swaps the cycles carried in packets with $y+1$, then sends out those packets at cycle $y+1$. This cycle mapping (e.g., $x \rightarrow y+1$) can be realized as an adjustment value, and it exists between any pair of neighbor nodes. With this mapping, the receiving node can easily figure out when the received packets should be sent out, the only requirement is to carry the cycle identifier of sending node in the packets.

In right part of Figure 4, Node A sends a packet with cycle identifier x in cycle x indicated by the identifier. After received by Node B, the cycle identifier x in the packet will be modified by the adjustment value to get a new cycle identifier $y+1$. Then the identifier $y+1$ will replace the original identifier x . Finally, the packet with the cycle identifier $y+1$ will be sent by Node B in cycle $y+1$ indicated by the new identifier.




Figure 4: CQF & SDF

3.1. Three Queues

In CQF each port needs to maintain 2 (or 3) queues: one is used to buffer newly received packets, another one is used to store the packets that are going to be sent out, one more queue may be needed to avoid output starvation [scheduled-queues]. In SDF, at least 3 cyclic queues are maintained for each port on a node. A cyclic queue corresponds to a cycle.

As Figure 5 illustrated, a node may receive packets sent at two different cycles from a single upstream node due to the absence of time synchronization. Following the cycle mapping (i.e., $x \rightarrow y+1$), packets that carry cycle identifier x should be sent out by Node B at cycle $y+1$, and packets that carry cycle identifier $x+1$ should be sent out by Node B at cycle $y+2$. Therefore, two queues are needed to store the newly received packets, as well as one queue to store the sending packets. In order to absorb more link delay variation (such as on radio interface), more queues may be necessary.




Figure 5: Three Queues in SDF

3.2. Cycle Mapping

When this packet is received by Node B, some methods are possible how the forwarding plane could operate. In one method, Node B has a mapping determined by the control plane. Packets from (the link from) Node A indicating cycle x are mapping into cycle y+1. This mapping is necessary, because all the packets from one cycle of the sending node need to get into one cycle of the receiving node. This is called "configured cycle mapping".

Instead of configuring an explicit cycle mapping such as cycle x -> cycle y+1, the receiving Node B could also have the intelligence in the forwarding plane to recognize the first packet from (the link from) Node A that has a new cycle x number, and map this cycle x to a cycle y after the current cycle. We call this option "self synchronized cycle mapping".

3.2.1. Cycle Identifier Carrying

In self synchronized cycle mapping, cycle identifier needs to be carried in the SDF packets, so that an appropriate queue can be selected accordingly. That means 2 bits are needed in the three queues model of SDF, in order to identify different cycles between a pair of neighboring nodes. There are several ways to carry this 2 bits cycle identifier. This document does not yet aim to propose one, but gives an (incomplete) list of ideas:

- o DSCP of IPv4 Header
- o Traffic Class of IPv6 Header
- o TC of MPLS Header (used to be EXP)
- o EtherType of Ethernet Header
- o IPv6 Extension Header

- o TLV of SRv6
- o TC of MPLS-SR Header (used to be EXP)
- o Three labels/adjacency SIDs for MPLS-SR

4. Scalable Resource Reservation

SDF must work with some resource reservation mechanisms, that can fulfill the role of the Scalable Resource Reservation (SRR). This resource reservation guarantees the necessary network resources (e.g., bandwidth) when deterministic flows are scheduled including the slots through which the traffic travels hop-by-hop. Network nodes have to record how many network resources are reserved for a specific flow from when it starts to when it ends (e.g., `<flow_identifier, reserved_resource, start_time, end_time>`). Maintaining per-flow resource reservation state may be acceptable to edge nodes, but un-acceptable to core nodes. [draft-ietf-detnet-architecture] pointed out that aggregation must be supported for scalability.

SRR aggregates per-flow resource reservation states in each time slot following the steps:

1. Dividing time into time slots. Then the per-flow resource reservation message can be expressed as `<flow_identifier, reserved_resource, start_time_slot, num_time_slot>` accordingly, where `flow_identifier` is the identifier of a deterministic flow, `reserved_resource` indicates how much resource is reserved, `start_time_slot` is the number of time slot from which resource reservation starts (e.g., the time slot that a new resource reservation request generates), `num_time_slot` indicates how many time slots the resource will be reserved. Note that time slot here is irrelevant to the cycle in SDF.
2. Edge node still maintains per-flow resource reservation states. While core node calculates and maintains the sum of `reserved_resources` (or remaining resources) of each time slot. That is a core node just needs to maintain a variable for each time slot. A core node can maintain K time slots' resource reservation states, i.e., the aggregated resource reservation window of a core node is K.
3. New resource reservation request succeed only if there are sufficient resources along the path. That is every related core node's remaining resource is no less than the amount of newly request resource. Otherwise, the resource reservation request failed. Resource is reserved in unit of time slot, and at most K

time slots. If a flow wants to consecutively reserve resources after the new resource reservation request expired, edge node/host can send renewal request. Similar to new resource reservation request, renewal request also needs to carry the flow identifier (the same identifier as the flow identifier carried by the new resource reservation request), the amount of reserved resource (no more than the previous request), as well as the number of time slot that the resource will be reserved. Edge node/host also can active teardown the resource reservation along the path.

4. After receiving the the per-flow resource reservation message, core nodes refresh their aggregated resource reservation windows accordingly. As item 2 specifies, core node may record the sum of reserved_resource or the remaining resource (remaining resource = capacity - sum of reserved_resource). If the sum of reserved resources is recorded, then core node should add the newly requested resource to the maintained resource in each related time slot. Otherwise if the remaining resource is recorded, then core node should subtract the newly requested resource to the maintained resource in each related time slot.

5. Performance Analysis

5.1. Queueing Delay

We consider forwarding from an LDN node A via an LDN node B to an LDN node C and call the single-hop LDN delay the time between a packet being sent by A and the time it is re-sent by B. This single-hop delay is composed from the A->B propagation delay and the single-hop queueing delay A->B.




Figure 6: Single-Hop Queueing Delay

As Figure 6 shows, cycle x of Node A will be mapped into cycle y+1 of Node B as long as the last packet sent from A->B is received within

the cycle y . If the last packet is re-sent out by B at the end of cycle $y+1$, then the largest single-hop queueing delay is $2*T$. Therefore the end-to-end queueing delay's upper bound is $2*T*H$, where H is the number of hops.

If A did not forward the LDN packet from a prior LDN forwarder but is the actual traffic source, then the packet may have been delayed by a gate function before it was sent to B. The delay of this function is outside of scope for the LDN delay considerations. If B is not forwarding the LDN packet but the final receiver, then the packet may not need to be queued and released in the same fashion to the receiver as it would be queued/released to a downstream LDN node, so if a path has one source followed by N LDN forwarders followed by one receivers, this should be considered to be a path with $N-1$ LDN hops for the purpose of latency and jitter calculations.

5.2. Jitter

Considering the simplest scenario one hop forwarding at first, suppose Node A is the upstream node of Node B, the packet sent from Node A at cycle x will be received by Node B at cycle y as Figure 7 shows.

- The best situation is Node A sends packet at the end of cycle x , and Node B receives packet at the beginning of cycle y , then the delay is denoted by w ;
- The worst situation is Node A sends packet at the beginning of cycle x , and Node B receives packet at the end of cycle y , then the delay = $w + \text{length of cycle } x + \text{length of cycle } y = w + 2*T$;
- Hence the jitter's upper bound of this simplest scenario = worst case - best case = $2*T$.

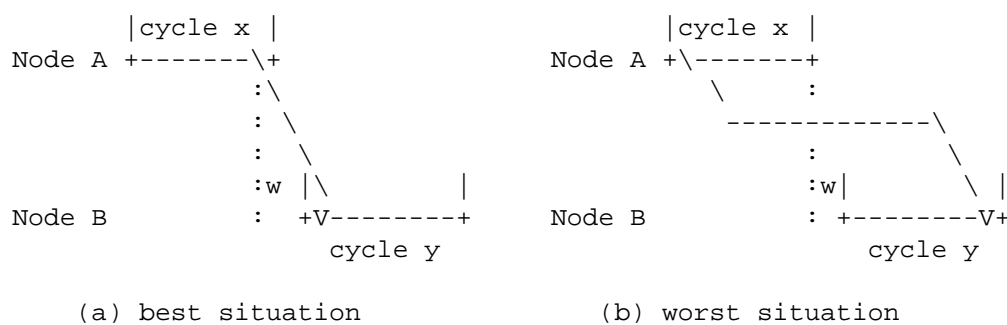




Figure 7: Jitter Analysis for One Hop Forwarding

Next considering two hops forwarding as Figure 8 shows.

- The best situation is Node A sends packet at the end of cycle x, and Node C receives packet at the beginning of cycle z, then the delay is denoted by w' ;
- The worst situation is Node A sends packet at the beginning of cycle x, and Node C receives packet at the end of cycle z, then the delay = $w' + \text{length of cycle x} + \text{length of cycle z} = w' + 2 * T$;
- Hence the jitter's upper bound = worst case - best case = $2 * T$.



(a) best situation



(b) worst situation

Figure 8: Jitter Analysis for Two Hops Forwarding

And so on. For multi-hop forwarding, the end-to-end delay will increase as the number of hops increases, while the delay variation (jitter) still does not exceed $2 * T$.

6. IANA Considerations

This document makes no request of IANA.

7. Security Considerations

Security issues have been carefully considered in [draft-ietf-detnet-security]. More discussion is TBD.

8. Acknowledgements

TBD.

9. Normative References

- [draft-ietf-detnet-architecture]
"DetNet Architecture", <<https://datatracker.ietf.org/doc/draft-ietf-detnet-architecture/>>.
- [draft-ietf-detnet-dp-sol]
"DetNet Data Plane Encapsulation",
<<https://datatracker.ietf.org/doc/draft-ietf-detnet-dp-sol/>>.
- [draft-ietf-detnet-problem-statement]
"DetNet Problem Statement",
<<https://datatracker.ietf.org/doc/draft-ietf-detnet-problem-statement/>>.
- [draft-ietf-detnet-security]
"DetNet Security Considerations",
<<https://datatracker.ietf.org/doc/draft-ietf-detnet-security/>>.
- [draft-ietf-detnet-use-cases]
"DetNet Use Cases", <<https://datatracker.ietf.org/doc/draft-ietf-detnet-use-cases/>>.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [scheduled-queues]
"Scheduled queues, UBS, CQF, and Input Gates",
<<http://www.ieee802.org/1/files/public/docs2015/new-nfinn-input-gates-0115-v04.pdf>>.

Authors' Addresses

Li Qiang (editor)
Huawei
Beijing
China

Email: qiangli3@huawei.com

Bingyang Liu
Huawei
Beijing
China

Email: liubingyang@huawei.com

Toerless Eckert (editor)
Huawei USA - Futurewei Technologies Inc.
2330 Central Expy
Santa Clara 95050
USA

Email: tte+ietf@cs.fau.de

Liang Geng
China Mobile
Beijing
China

Email: gengliang@chinamobile.com

Lei Wang
China Mobile
Beijing
China

Email: wangleiyjy@chinamobile.com